

**Structure Without Weights: Recovering Document Layout
from the Text Layer Alone**

Ada Lovelace

Analytical Engine Laboratory

Author Note

- ¹ Correspondence concerning this article should be addressed to Ada Lovelace, Analytical
² Engine Laboratory, London. This manuscript is a synthetic fixture and describes no real study.

Introduction

A converter that joins lines with a space will emit a broken word at every such break. Tables were drawn with visible ruling lines so that the vector-based detector fires before the text-alignment fallback. Footnotes were set two points smaller than the body and anchored in the bottom fifth of the column. Their labels are superscript digits, and matching digits appear in the body at the point of reference.

The abstract runs across the full measure while the body is set in two columns, which is the arrangement most journals use. The manuscripts we care about in practice are less tidy than this, and we return to them in the discussion. The text-layer path handles digital publications; a raster copy of the same file drives the recognition path. Both paths converge on the same block structure before any processor runs.

Nothing in the pipeline depends on a downloaded model, which is the constraint that motivated the project. Line heights are clustered to recover heading levels when a document has no section numbers. When numbers are present they win, because a numbered heading states its own depth. Captions are recognised by their leading label and attached to the nearest figure or table above them.

Theory and Hypotheses

Equations are left as images for a later pass rather than transcribed into notation that would be wrong half the time. The remaining processors are direct ports and are documented against the source files they come from. Every threshold in the code was set by looking at a failure, not by tuning against a corpus. That makes the thresholds easy to defend and easy to revise when a new failure appears.

We report where the approach breaks so that a reader can decide whether it fits their documents. Multi-line table cells remain the weakest point and are listed as partial support. Journals that style every heading identically defeat the height clustering and lose a level. Scanned pages inherit every error of the recogniser, including stray page numbers in the middle of a paragraph.

Hypothesis 1: Documents converted with stream order will show fewer reading-order errors than documents converted with geometric block sorting.

None of these failures corrupts the surrounding text; they degrade the structure rather than the content. The synthetic set is regenerated from scripts, so a fixture can be changed and the change reviewed. Expected outputs are committed beside the fixtures and compared byte for byte. A behaviour change that is intended is recorded by rewriting the expected files in the same commit.

Method

35 An unintended change fails the build before it reaches a release. The vendored table code
36 is treated as third-party and is not edited locally. Where the original relied on a learned layout
37 model, a font-size rule stands in for it. The rule is coarse, but it is inspectable, and a wrong
38 answer can be traced to a number in the source.

39 Reviewers asked whether the approach generalises beyond journal articles; we make no
40 such claim. Books, slides and forms have layouts these heuristics were never shown and would
41 misread. The scope is deliberately narrow: articles with a text layer, converted for reading by a
42 language model. Within that scope the results are stable across the publishers we tried.

43 Outside it, users should expect to check the output before trusting it. A second reviewer
44 asked for timings; conversion runs at a few pages per second on a laptop. Most of that time is
45 spent in the table detector, which is invoked on every page. Skipping it on pages with no ruling
46 lines would halve the runtime at no cost to accuracy.